

- abs/2412.13459 (2024).
- [61] Junda He, Christoph Treude, and David Lo. 2025. LLM-Based Multi-Agent Systems for Software Engineering: Literature Review, Vision, and the Road Ahead. *TOSEM* 34, 5 (2025), 1–30.
 - [62] Runzhi He, Hao He, Yuxia Zhang, and Minghui Zhou. 2023. Automating Dependency Updates in Practice: An Exploratory Study on GitHub Dependabot. *TSE* 49, 8 (2023), 4004–4022.
 - [63] Manuel Hoffmann, Sam Boysel, Frank Nagle, Sida Peng, and Kevin Xu. 2024. *Generative AI and the Nature of Work*. Technical Report. CESifo Working Paper.
 - [64] Xinyi Hou, Yanjie Zhao, Yue Liu, Zhou Yang, Kailong Wang, Li Li, Xiapu Luo, David Lo, John Grundy, and Haoyu Wang. 2024. Large Language Models for Software Engineering: A Systematic Literature Review. *TOSEM* 33, 8 (2024), 220:1–220:79.
 - [65] Rasha Ahmad Husein, Hala Aburajouh, and Cagatay Catal. 2025. Large language models for code completion: A systematic literature review. *Comput. Stand. Interfaces* 92 (2025), 103917.
 - [66] Saki Imai. 2022. Is GitHub Copilot a Substitute for Human Pair-programming? An Empirical Study. In *ICSE*. ACM, 319–321.
 - [67] Haolin Jin, Linghan Huang, Haipeng Cai, Jun Yan, Bo Li, and Huaming Chen. 2024. From LLMs to LLM-based Agents for Software Engineering: A Survey of Current, Challenges and Future. *CoRR* abs/2408.02479 (2024).
 - [68] Eirini Kalliamvakou, Georgios Gousios, Kelly Blincoe, Leif Singer, Daniel M. Germán, and Daniela E. Damian. 2016. An in-depth study of the promises and perils of mining GitHub. *EMSE* 21, 5 (2016), 2035–2071.
 - [69] Mohammed Kharmah, Soohyeon Choi, Mohammed AlKhanafseh, and David Mohaisen. 2025. Security and Quality in LLM-Generated Code: A Multi-Language, Multi-Model Analysis. *CoRR* abs/2502.01853 (2025).
 - [70] Raphaël Khoury, Anderson R. Avila, Jacob Brunelle, and Baba Mamadou Camara. 2023. How Secure is Code Generated by ChatGPT?. In *SMC*. IEEE, 2445–2451.
 - [71] Barbara A. Kitchenham and Shari Lawrence Pfleeger. 1996. Software Quality: The Elusive Target. *IEEE Softw.* 13, 1 (1996), 12–21.
 - [72] Philippe Kruchten, Robert L Nord, and Ipek Ozkaya. 2012. Technical debt: From metaphor to theory and practice. *IEEE Softw.* 29, 6 (2012), 18–21.
 - [73] Anand Kumar, Vishal Khare, Deepak Sharma, Satyam Kumar, Vijay Saini, Anshul Yadav, Sachendra Jain, Ankit Rana, Pratham Verma, Vaibhav Meena, et al. 2025. Intuition to Evidence: Measuring AI's True Impact on Developer Productivity. *CoRR* abs/2509.19708 (2025).
 - [74] M. J. Lawrence. 1981. Programming methodology, organizational environment, and programming productivity. *JSS* 2, 3 (1981), 257–269.
 - [75] Shuang Li, Yuntao Cheng, Jinfu Chen, Jifeng Xuan, Sen He, and Weiyi Shang. 2024. Assessing the Performance of AI-Generated Code: A Case Study on GitHub Copilot. In *ISSRE*. IEEE, 216–227.
 - [76] Jenny T. Liang, Chenyang Yang, and Brad A. Myers. 2024. A Large-Scale Survey on the Usability of AI Programming Assistants: Successes and Challenges. In *ICSE*. ACM, 52:1–52:13.
 - [77] Junwei Liu, Kaixin Wang, Yixuan Chen, Xin Peng, Zhenpeng Chen, Lingming Zhang, and Yiling Lou. 2024. Large Language Model-Based Agents for Software Engineering: A Survey. *CoRR* abs/2409.02977 (2024).
 - [78] Zhijie Liu, Yutian Tang, Xiapu Luo, Yuming Zhou, and Liang Feng Zhang. 2024. No Need to Lift a Finger Anymore? Assessing the Quality of Code Generation by ChatGPT. *TSE* 50, 6 (2024), 1548–1584.
 - [79] Junyi Lu, Lei Yu, Xiaojia Li, Li Yang, and Chun Zuo. 2023. LLaMA-Reviewer: Advancing Code Review Automation with Large Language Models through Parameter-Efficient Fine-Tuning. In *ISSRE*. IEEE, 647–658.
 - [80] Robert C Martin. 2009. *Clean Code: A Handbook of Agile Software Craftsmanship*. Pearson Education.
 - [81] Boris Martinović and Robert Rozić. 2024. Impact of AI tools on software development code quality. In *ICDTEAI*. Springer, 241–256.
 - [82] Thomas J. McCabe. 1976. A Complexity Measure. *TSE* 2, 4 (1976), 308–320.
 - [83] Audris Mockus, Roy T. Fielding, and James D. Herbsleb. 2002. Two case studies of open source software development: Apache and Mozilla. *TOSEM* 11, 3 (2002), 309–346.
 - [84] Nuthan Munaiah, Steven Kroh, Craig Cabrey, and Meiyappan Nagappan. 2017. Curating GitHub for engineered software projects. *EMSE* 22, 6 (2017), 3219–3253.
 - [85] Emerson R. Murphy-Hill, Ciera Jaspán, Caitlin Sadowski, David C. Shepherd, Michael Phillips, Collin Winter, Andrea Knight, Edward K. Smith, and Matthew Jorde. 2021. What Predicts Software Developers' Productivity? *TSE* 47, 3 (2021), 582–594.
 - [86] Keitaro Nakasai, Hideaki Hata, and Kenichi Matsumoto. 2018. Are donation badges appealing?: A case study of developer responses to Eclipse bug reports. *IEEE Software* 36, 3 (2018), 22–27.
 - [87] Nhan Nguyen and Sarah Nadi. 2022. An Empirical Evaluation of GitHub Copilot's Code Suggestions. In *MSR*. ACM, 1–5.
 - [88] Sanghak Oh, Kiho Lee, Seonhye Park, Doowon Kim, and Hyoungshick Kim. 2024. Poisoned ChatGPT Finds Work for Idle Hands: Exploring Developers' Coding Practices with Insecure Suggestions from Poisoned AI Models. In *S&P*. IEEE, 1141–1159.
 - [89] Edson Oliveira, Eduardo Fernandes, Igor Steinmacher, Marco Cristo, Tayana Conte, and Alessandro Garcia. 2020. Code and commit metrics of developer productivity: A study on team leaders perceptions. *EMSE* 25, 4 (2020), 2519–2549.
 - [90] Leon J. Osterweil. 1996. Strategic Directions in Software Quality. *ACM Comput. Surv.* 28, 4 (1996), 738–750.
 - [91] Ruchika Pandey, Prabhat Singh, Raymond Wei, and Shaila Shankar. 2024. Transforming Software Development: Evaluating the Efficiency and Challenges of GitHub Copilot in Real-World Projects. *CoRR* abs/2406.17910 (2024).
 - [92] Elise Paradis, Kate Grey, Quinn Madison, Daye Nam, Andrew Macvean, Vahid Meimand, Nan Zhang, Ben Ferrari-Church, and Satish Chandra. 2025. How much does AI impact development speed? An enterprise-based randomized controlled trial. In *ICSE-SEIP*. IEEE, 618–629.
 - [93] Hammond Pearce, Baleegh Ahmad, Benjamin Tan, Brendan Dolan-Gavitt, and Ramesh Karri. 2022. Asleep at the Keyboard? Assessing the Security of GitHub Copilot's Code Contributions. In *S&P*. IEEE, 754–768.
 - [94] Sida Peng, Eirini Kalliamvakou, Peter Cihon, and Mert Demirel. 2023. The Impact of AI on Developer Productivity: Evidence from GitHub Copilot. *CoRR* abs/2302.06590 (2023).
 - [95] Neil Perry, Megha Srivastava, Deepak Kumar, and Dan Boneh. 2023. Do Users Write More Insecure Code with AI Assistants?. In *CCS*. ACM, 2785–2799.
 - [96] Jeremy A Rassen, Abhi A Shelat, Jessica Myers, Robert J Glynn, Kenneth J Rothman, and Sebastian Schneeweiss. 2012. One-to-many propensity score matching in cohort studies. *Pharmacoepidemiol. Drug Saf.* 21 (2012), 69–80.
 - [97] Gustavo Sandoval, Hammond Pearce, Teo Nys, Ramesh Karri, Siddharth Garg, and Brendan Dolan-Gavitt. 2023. Lost at C: A User Study on the Security Implications of Large Language Model Code Assistants. In *USENIX Security*. USENIX, 2205–2222.
 - [98] Suproteem K Sarkar. 2025. AI agents, productivity, and higher-order thinking: Early evidence from software development. *Available at SSRN* 5713646 (2025).
 - [99] Max Schäfer, Sarah Nadi, Aryaz Eghbali, and Frank Tip. 2024. An Empirical Evaluation of Using Large Language Models for Automated Unit Test Generation. *TSE* 50, 1 (2024), 85–105.
 - [100] Ingo Scholtes, Pavlin Mavrodiev, and Frank Schweitzer. 2016. From Aristotle to Ringelmann: a large-scale analysis of team productivity and coordination in Open Source Software projects. *EMSE* 21, 2 (2016), 642–683.
 - [101] Mohammed Latif Siddiq, Shafayat H. Majumder, Maisha R. Mim, Sourov Jadodia, and Joanna C. S. Santos. 2022. An Empirical Study of Code Smells in Transformer-based Code Generation Techniques. In *SCAM*. IEEE, 71–82.
 - [102] Fangchen Song, Ashish Agarwal, and Wen Wen. 2024. The Impact of Generative AI on Collaborative Open-Source Software Development: Evidence from GitHub Copilot. *CoRR* abs/2410.02091 (2024).
 - [103] César Soto-Valero, Thomas Durieux, and Benoit Baudry. 2021. A longitudinal analysis of bloated Java dependencies. In *FSE*. ACM, 1021–1031.
 - [104] Margaret-Anne Storey, Brian Houck, and Thomas Zimmermann. 2022. How developers and managers define and trade productivity for quality. In *CHASE*. 26–35.
 - [105] Viktoria Stray, Elias Goldmann Brandtzaeg, Viggo Tellefsen Wivestad, Astri Barbala, and Nils Brede Moe. 2025. Developer Productivity With and Without GitHub Copilot: A Longitudinal Mixed-Methods Case Study. *CoRR* abs/2509.20353 (2025).
 - [106] Elizabeth A Stuart, Brian K Lee, and Finbarr P Leacy. 2013. Prognostic score-based balance measures can be a useful diagnostic for propensity score methods in comparative effectiveness research. *Journal of Clinical Epidemiology* 66, 8 (2013), S84–S90.
 - [107] Lukas Twist, Jie M. Zhang, Mark Harman, Don Syme, Joost Noppen, and Detlef D. Nauck. 2025. LLMs Love Python: A Study of LLMs' Bias for Programming Languages and Libraries. *CoRR* abs/2503.17181 (2025).
 - [108] Priyan Vaithilingam, Tianyi Zhang, and Elena L. Glassman. 2022. Expectation vs. Experience: Evaluating the Usability of Code Generation Tools Powered by Large Language Models. In *CHI Extended Abstracts*. ACM, 332:1–332:7.
 - [109] Chong Wang, Kaifeng Huang, Jian Zhang, Yebo Feng, Luyue Zhang, Yang Liu, and Xin Peng. 2025. LLMs Meet Library Evolution: Evaluating Deprecated API Usage in LLM-based Code Completion. In *ICSE*. IEEE, 781–781.
 - [110] Miku Watanabe, Hao Li, Yutaro Kashiwa, Brittany Reid, Hajimu Iida, and Ahmed E Hassan. 2025. On the use of agentic coding: An empirical study of pull requests on GitHub. *CoRR* abs/2509.14745 (2025).
 - [111] Doron Yeverechyahu, Raveesh Mayya, and Gal Oestreicher-Singer. 2024. The Impact of Large Language Models on Open-source Innovation: Evidence from GitHub Copilot. In *ICIS*. AIS.
 - [112] Jiasheng Zheng, Boxi Cao, Zhengzhao Ma, Ruotong Pan, Hongyu Lin, Yaojie Lu, Xianpei Han, and Le Sun. 2024. Beyond correctness: Benchmarking multi-dimensional code generation for large language models. *CoRR* abs/2407.11470 (2024).
 - [113] Albert Ziegler, Eirini Kalliamvakou, X. Alice Li, Andrew Rice, Devon Rifkin, Shawn Simister, Ganesh Sittampalam, and Edward Aftandilian. 2022. Productivity assessment of neural code completion. In *MAPS@PLDI*. ACM, 21–29.